

1a) pandas library funs

```
import pandas as pd

# Create a sample DataFrame
data = {
    'Name': ['John', 'Jane', 'Bob', 'Alice'],
    'Age': [25, 30, 22, 28],
    'City': ['New York', 'San Francisco', 'Los Angeles', 'Chicago']
}

df = pd.DataFrame(data)

# Export data to a CSV file
csv_filename = 'sample_data.csv'
df.to_csv(csv_filename, index=False)
print(f'Data exported to {csv_filename}')

# Import data from the CSV file
imported_df = pd.read_csv(csv_filename)

# Display the imported DataFrame
print("\nImported DataFrame:")
print(imported_df)

# Export data to an Excel file
excel_filename = 'sample_data.xlsx'
df.to_excel(excel_filename, index=False, sheet_name='Sheet1')
print(f'\nData exported to {excel_filename}')

# Import data from the Excel file
imported_df_excel = pd.read_excel(excel_filename, sheet_name='Sheet1')

# Display the imported DataFrame from Excel
print("\nImported DataFrame from Excel:")
print(imported_df_excel)
```

1B.& 6A.) Logistic Regression model

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Load the Iris dataset
from sklearn.datasets import load_iris
iris_data = load_iris()
iris_df = pd.DataFrame(iris_data.data, columns=iris_data.feature_names)
iris_df['target'] = iris_data.target

# Consider a binary classification problem (0 or 1)
iris_df['binary_target'] = (iris_df['target'] == 0).astype(int)

# Split the dataset into features (X) and binary target variable (y)
X = iris_df.drop(['target', 'binary_target'], axis=1)
y = iris_df['binary_target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the logistic regression model
log_reg = LogisticRegression(random_state=42)

# Train the model
log_reg.fit(X_train, y_train)

# Make predictions on the test set
y_pred = log_reg.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_rep)
print("\nConfusion Matrix:\n", conf_matrix)

# Plot the confusion matrix
plt.figure(figsize=(6, 4))
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Class 0',
'Class 0'], yticklabels=['Not Class 0', 'Class 0'])

plt.title('Confusion Matrix')

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.show()

# ROC Curve

from sklearn.metrics import roc_curve, auc, confusion_matrix

import matplotlib.pyplot as plt

# Get the predicted probabilities for class 1 (positive class)
y_pred_proba = log_reg.predict_proba(X_test)[: , 1]

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculate the Area Under the ROC Curve (AUC)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))

plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)

plt.plot([0, 1], [0, 1], color='red', lw=2, linestyle='--')

plt.xlim([0.0, 1.0])

plt.ylim([0.0, 1.05])

plt.xlabel('False Positive Rate')

plt.ylabel('True Positive Rate')

plt.title('Receiver Operating Characteristic (ROC) Curve')

plt.legend(loc="lower right")

plt.show()
```

2A&3A&13B) Data pre-processing (EDA)

#import necessary libraries

import pandas as pd

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

Load the dataset

used_cars_df = pd.read_csv('used_cars.csv')

Display the first few rows of the dataset

print(used_cars_df.head())

#Get the basic information about the dataset

print(used_cars_df.info())

Summary statistics of the numerical columns

print(used_cars_df.describe())

#Check for missing values

print(used_cars_df.isnull().sum())

#Check the distribution of the categorical variables

for column in used_cars_df.select_dtypes(include='object').columns:

print(used_cars_df[column].value_counts())

#Visualize the distribution of numerical features

sns.pairplot(used_cars_df)

plt.show()

#Correlation Matrix

correlation_matrix = used_cars_df.corr()

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')

plt.title('Correlation Matrix')

plt.show()

#Visualize the distribution of categorical variables

for column in used_cars_df.select_dtypes(include='object').columns:

sns.countplot(x=column, data=used_cars_df)

plt.title(f'Distribution of {column}')

```
plt.xticks(rotation=45)
```

```
plt.show()
```

3B&11B) Naïve byes theorem

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Sample text data
data = {'text': ["I love programming", "Machine learning is fascinating", "Spam emails are annoying",
"Python is a great language", "Buy our new product now"]}
labels = [1, 1, 0, 1, 0] # 1 for positive, 0 for negative

# Create a DataFrame
df = pd.DataFrame({'text': data['text'], 'label': labels})

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'], test_size=0.2,
random_state=42)

# Vectorize the text data using CountVectorizer
vectorizer = CountVectorizer()
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Initialize the Naive Bayes model (Multinomial Naive Bayes for text data)
naive_bayes = MultinomialNB()

# Train the model
naive_bayes.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = naive_bayes.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_rep)
print("\nConfusion Matrix:\n", conf_matrix)
```

3B&11B DATA VISUALIZATION

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

iris = sns.load_dataset('Iris')
plt.figure(figsize=(10,6))
sns.scatterplot(x='sepal_length', y='sepal_width', hue='species', data=iris)
plt.title('Scatter Plot of Sepal length Vs Sepal Width')
plt.show()

sns.pairplot(iris, hue='species', height=2.5)
plt.subtitle('Pairwise Relationships and Distributions')
plt.show()

plt.figure(figsize=(15,8))
for i, feature in enumerate(iris.columns[:-1]):
    plt.subplot(2,2, i+1)
    sns.boxplot(x='species', y=feature, data=iris)
    plt.title(f'Boxplot of {feature}')
plt.tight_layout()
plt.show()

plt.figure(figsize=(15,8))
for i, feature in enumerate(iris.columns[:-1]):
    plt.subplot(2,2, i+1)
    sns.violinplot(x='species', y=feature, data=iris)
    plt.title(f'Violinplot of {feature}')
plt.tight_layout()
plt.show()

correlation_matrix = iris.corr()
```

```

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()

plt.figure(figsize=(10,6))
sns.histplot(iris['sepal_length'], kde=True)
plt.title('Distribution of Sepal Length')
plt.show()

plt.figure(figsize=(8,6))
sns.countplot(x='species', data=iris)
plt.title('Count Plot of Species')
plt.show()

```

4A) DECISION TREE CLASS

```

import pandas as pd
import numpy as np

df=pd.read_csv('PlayTennis.csv')

from sklearn.preprocessing import LabelEncoder

enc = LabelEncoder()

df_num_cat=pd.DataFrame()

df_num_cat['Outlook']= enc.fit_transform(df['Outlook'])

df_num_cat['Temperature']= enc.fit_transform(df['Temperature'])

df_num_cat['Humidity']= enc.fit_transform(df['Humidity'])

df_num_cat['Wind']= enc.fit_transform(df['Wind'])

df_num_cat['Play Tennis']= enc.fit_transform(df['Play Tennis'])

df_num_cat

x=df_num_cat.drop(['Play Tennis'],axis=1)

y=df_num_cat['Play Tennis']

# Split dataset into training set and test set

from sklearn.model_selection import train_test_split # Import train_test_split function

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=1) #
70% training

and 30% test

```

```

print(X_train)
print(X_test)
print(y_train)
print(y_test)
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
df_clf=DecisionTreeClassifier(criterion='entropy')
#df_clf=DecisionTreeClassifier()
df_clf.fit(X_train,y_train)
y_pred=df_clf.predict(X_test)
print(y_train)
print(y_test)
print(y_pred)
print("Accuracy")
print(metrics.accuracy_score(y_test,y_pred))
# Plot the decision tree
plt.figure(figsize=(12, 8))
plot_tree(df_clf, filled=True, feature_names=X_train.columns, class_names=df['Play
Tennis'].unique())
plt.show()

```

4B) NUMPY PROG

```

import numpy as np
# Define two vectors
vector1 = np.array([1, 2, 3])
vector2 = np.array([4, 5, 6])
# Element-wise multiplication
result = vector1 * vector2
# Display the result
print("Vector 1:", vector1)
print("Vector 2:", vector2)
print("Element-wise Multiplication:", result)

```


5A)SIMPLE & MULTI LINEARS

```
import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

# Sample data: one feature (e.g., study hours) and one target (e.g., marks)
X = np.array([[1], [2], [3], [4], [5]]) # Feature
y = np.array([2, 4, 5, 4, 5])          # Target

# Create and train the model
model = LinearRegression()
model.fit(X, y)

# Predict
y_pred = model.predict(X)

# Print coefficients
print("Simple Linear Regression:")
print("Intercept:", model.intercept_)
print("Slope:", model.coef_)

# Plotting
plt.scatter(X, y, color='blue')
plt.plot(X, y_pred, color='red')
plt.title('Simple Linear Regression')
plt.xlabel('X')
plt.ylabel('y')
plt.show()
```

MULTIPLE :

```
import numpy as np
from sklearn.linear_model import LinearRegression

# Sample data: multiple features (e.g., study hours and sleep hours)
X = np.array([[1, 7], [2, 8], [3, 8], [4, 6], [5, 5]]) # Features
y = np.array([50, 60, 65, 70, 75])                    # Target

# Create and train the model
model = LinearRegression()
```

```

model.fit(X, y)

# Predict
y_pred = model.predict(X)

# Print coefficients
print("\nMultiple Linear Regression:")
print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)

# Display predicted values
print("Predicted y:", y_pred)

```

5B) PANDAS PROG TO ADD,SUBTRACT,MULTIPLY AND DIVIDE 2 PANDAS SERIES

```

import pandas as pd

ds1 = pd.Series([2, 4, 6, 8, 10])
ds2 = pd.Series([1, 3, 5, 7, 9]) ds = ds1 + ds2
print("Add two Series:") print(ds)

print("Subtract two Series:") ds = ds1 - ds2
print(ds)

print("Multiply two Series:") ds = ds1 * ds2
print(ds)

print("Divide Series1 by Series2:") ds = ds1 / ds2
print(ds)

```

2B&7A NAÏVE BAYES

```

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Sample text data
data = {'text': ["I love programming", "Machine learning is fascinating", "Spam emails are annoying",
"Python is a great language", "Buy our new product now"]}

labels = [1, 1, 0, 1, 0] # 1 for positive, 0 for negative

```

```

# Create a DataFrame
df = pd.DataFrame({'text': data['text'], 'label': labels})

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'], test_size=0.2,
random_state=42)

# Vectorize the text data using CountVectorizer
vectorizer = CountVectorizer()
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Initialize the Naive Bayes model (Multinomial Naive Bayes for text data)
naive_bayes = MultinomialNB()

# Train the model
naive_bayes.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = naive_bayes.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_rep)
print("\nConfusion Matrix:\n", conf_matrix)

```

7B) FRIDAY

```

# Given probabilities
P_Friday_and_Absent = 0.03
P_Friday = 0.20

# Bayes' Theorem:  $P(\text{Absent} \mid \text{Friday}) = P(\text{Friday and Absent}) / P(\text{Friday})$ 
P_Absent_given_Friday = P_Friday_and_Absent / P_Friday

# Print the result
print("Probability that a student is absent given that today is Friday:")
print(f"P(Absent | Friday) = {P_Absent_given_Friday:.2f}")

```

8A) K-NEAREST NEIGHBOURS

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load the Iris dataset

from sklearn.datasets import load_iris

iris_data = load_iris()

iris_df = pd.DataFrame(iris_data.data, columns=iris_data.feature_names)

iris_df['target'] = iris_data.target

# Split the dataset into features (X) and target variable (y)

X = iris_df.drop('target', axis=1)

y = iris_df['target']

# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the K-Nearest Neighbors classifier

knn_classifier = KNeighborsClassifier(n_neighbors=3)

# Train the model

knn_classifier.fit(X_train, y_train)

# Make predictions on the test set

y_pred = knn_classifier.predict(X_test)

# Evaluate the model

accuracy = accuracy_score(y_test, y_pred)

classification_rep = classification_report(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

print("Accuracy:", accuracy)

print("\nClassification Report:\n", classification_rep)

print("\nConfusion Matrix:\n", conf_matrix)
```

8B) RESHAPING

```
import numpy as np

import pandas as pd
```

```

# -----
# (i) Reshaping the Data
# -----
# Create a 1D NumPy array
data = np.array([1, 2, 3, 4, 5, 6])
# Reshape it into 2D (2 rows, 3 columns)
reshaped_data = data.reshape((2, 3))
print("Original Data (1D):")
print(data)
print("\nReshaped Data (2D):")
print(reshaped_data)
# -----
# (ii) Filtering the Data
# -----
# Create a Pandas DataFrame for filtering
df = pd.DataFrame({
    'Student': ['Alice', 'Bob', 'Charlie', 'David', 'Eva'],
    'Marks': [85, 40, 95, 35, 75]
})
# Filter students who scored more than 50
filtered_df = df[df['Marks'] > 50]
print("\nOriginal DataFrame:")
print(df)
print("\nFiltered Data (Marks > 50):")
print(filtered_df)
9A&9B ) K- MEANS
import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt

```

```

# Create a synthetic dataset
data, true_labels = make_blobs(n_samples=300, centers=4, random_state=42)

# Convert the data to a DataFrame
df = pd.DataFrame(data, columns=['Feature1', 'Feature2'])

# Visualize the original data
plt.scatter(df['Feature1'], df['Feature2'], s=50)
plt.title('Original Data')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()

# Initialize the K-Means model
kmeans = KMeans(n_clusters=4, random_state=42)

# Fit the model to the data
kmeans.fit(df)

# Get the cluster labels and centroids
cluster_labels = kmeans.labels_
centroids = kmeans.cluster_centers_

# Add cluster labels to the DataFrame
df['Cluster'] = cluster_labels

# Visualize the clustered data
plt.scatter(df['Feature1'], df['Feature2'], c=df['Cluster'], cmap='viridis', s=50)
plt.scatter(centroids[:, 0], centroids[:, 1], marker='X', s=200, color='red', label='Centroids')
plt.title('K-Means Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.show()

```

10A&16B) PCA

```

import pandas as pd

from sklearn.datasets import load_iris
from sklearn.decomposition import PCA

```

```
import matplotlib.pyplot as plt

# Load the Iris dataset
iris_data = load_iris()

X = iris_data.data
y = iris_data.target

class_names = iris_data.target

# Convert the data to a DataFrame
df = pd.DataFrame(X, columns=iris_data.feature_names)

# Standardize the data (optional but recommended for PCA)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

X_standardized = scaler.fit_transform(X)

# Initialize PCA with the number of components to retain
n_components = 2 # Set the desired number of components
pca = PCA(n_components=n_components)

# Fit and transform the data
X_pca = pca.fit_transform(X_standardized)

# Create a DataFrame from the PCA-transformed data
df_pca = pd.DataFrame(X_pca, columns=[f'PC{i+1}' for i in range(n_components)])
df_pca['Target'] = y

# Visualize the data before and after PCA
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Before PCA
plt.subplot(1, 2, 2)

for i in range(3):
    plt.scatter(X_pca[y == i, 0], X_pca[y == i, 1], c=[plt.cm.viridis(i / 2.0)], edgecolors='k', s=50,
        label=class_names[i])
ax1.scatter(X[:, 0], X[:, 1], c=y, cmap='viridis', edgecolors='k', s=50)
ax1.set_title('Original Data')
ax1.set_xlabel('Feature 1')
ax1.set_ylabel('Feature 2')
```

```

plt.legend()

# After PCA

ax2.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolors='k', s=50)

ax2.set_title('Data After PCA')

ax2.set_xlabel('Principal Component 1')

ax2.set_ylabel('Principal Component 2')

plt.legend()

plt.show()

```

10B) MEAN,MEDIAN,MODE AND VARIANCE

```

import pandas as pd

from scipy import stats

# Sample dataset (you can also load from CSV using pd.read_csv)

data = {

    'Scores': [88, 92, 79, 93, 85, 91, 76, 89, 95, 92]

}

# Create a DataFrame

df = pd.DataFrame(data)

# Compute statistical measures

mean_val = df['Scores'].mean()

median_val = df['Scores'].median()

mode_val = df['Scores'].mode()[0]    # mode() returns a Series

variance_val = df['Scores'].var()

std_dev_val = df['Scores'].std()

# Display the results

print("Descriptive Statistics:")

print(f"Mean: {mean_val}")

print(f"Median: {median_val}")

print(f"Mode: {mode_val}")

print(f"Variance: {variance_val}")

print(f"Standard Deviation: {std_dev_val}")

```


11A) LDA

```
from sklearn.datasets import load_iris

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

import matplotlib.pyplot as plt

# Load the iris dataset

iris = load_iris()

X = iris.data

y = iris.target

class_names = iris.target_names

# Perform Linear Discriminant Analysis (LDA)

lda = LinearDiscriminantAnalysis(n_components=2)

X_lda = lda.fit_transform(X, y)

# Plot the results

plt.figure(figsize=(12, 5))

# Before LDA

plt.subplot(1, 2, 1)

for i in range(3):

    plt.scatter(X[y == i, 0], X[y == i, 1], c=[plt.cm.viridis(i / 2.0)], edgecolors='k', s=50,

    label=class_names[i])

plt.title('Original Data')

plt.xlabel('Feature 1')

plt.ylabel('Feature 2')

plt.legend()

# After LDA

plt.subplot(1, 2, 2)

for i in range(3):

    plt.scatter(X_lda[y == i, 0], X_lda[y == i, 1], c=[plt.cm.viridis(i / 2.0)], edgecolors='k', s=50,

    label=class_names[i])

plt.title('Data After LDA')

plt.xlabel('Linear Discriminant 1')

plt.ylabel('Linear Discriminant 2')
```

```
plt.legend()
plt.tight_layout()
plt.show()
```

12A&14B) AND LOGIC GATE

```
import numpy as np

# Define input features:
input_features = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
print("Input features shape:", input_features.shape)

# Define target output:
target_output = np.array([0, 0, 0, 1])
print("Target output shape:", target_output.shape)
print("Target output:", target_output)

# Define weights:
weights = np.array([0.1, 0.2])
print("Initial weights:", weights)

# Bias weight:
bias = 0.3

# Learning Rate:
lr = 0.05

# Step function:
def step_function(x):
    return 1 if x >= 0 else 0

# Main logic for Perceptron:
# Running our code 10000 times:
for epoch in range(10000):
    total_error = 0
    for input_data, label in zip(input_features, target_output):
        inputs = input_data
        # Feedforward input:
        in_o = np.dot(inputs, weights) + bias
        # Feedforward output:
```

```

out_o = step_function(in_o)
# Calculating error
error = label - out_o
total_error += abs(error)
# Updating the weights values:
weights += lr * error * inputs
# Updating the bias weight value:
bias += lr * error
# Early stopping if total_error is 0
if total_error == 0:
    break
# Check the final values for weight and bias
print("Final weights:", weights)
print("Final bias:", bias)
# Test the trained perceptron
print("Testing the trained perceptron:")
for inputs, label in zip(input_features, target_output):
    prediction = step_function(np.dot(inputs, weights) + bias)
    print(f"Input: {inputs}, Predicted Output: {prediction}, Target Output: {label}")

```

12B) Write a python program to compute (i) Feature Normalization: Min-max normalization (ii) Filtering the data

```

import pandas as pd
from sklearn.preprocessing import MinMaxScaler
# Sample dataset
data = {
    'Student': ['A', 'B', 'C', 'D', 'E'],
    'Maths': [45, 89, 72, 60, 95],
    'Science': [38, 85, 66, 55, 90]
}
# Create DataFrame
df = pd.DataFrame(data)

```

```

print("Original Data:")
print(df)
# -----
# (i) Feature Normalization (Min-Max)
# -----
scaler = MinMaxScaler()
# Normalize only numeric columns
df[['Maths_normalized', 'Science_normalized']] = scaler.fit_transform(df[['Maths',
'Science']])
print("\nData after Min-Max Normalization:")
print(df)
# -----
# (ii) Filtering the Data
# -----
# Filter students who scored more than 80 in Maths
filtered_df = df[df['Maths'] > 80]
print("\nFiltered Data (Maths > 80):")
print(filtered_df)

```

6. 13 A) OR logic gate with 2-bit Binary input

```

import numpy as np
# Step activation function
def step_function(x):
    return 1 if x >= 0 else 0
# Perceptron training function
def train_perceptron(X, y, lr=0.1, epochs=10):
    weights = np.zeros(X.shape[1]) # 2 inputs
    bias = 0
    for epoch in range(epochs):
        print(f"Epoch {epoch + 1}")
        for i in range(len(X)):
            # Linear combination

```

```

    z = np.dot(X[i], weights) + bias
    # Activation
    y_pred = step_function(z)
    # Update weights and bias if prediction is wrong
    error = y[i] - y_pred
    weights += lr * error * X[i]
    bias += lr * error
    print(f"Input: {X[i]}, Target: {y[i]}, Predicted: {y_pred}, Error: {error}")
    print(f"Weights: {weights}, Bias: {bias}\n")
    return weights, bias
# Prediction function
def predict(X, weights, bias):
    return [step_function(np.dot(x, weights) + bias) for x in X]
# OR Gate dataset
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([0, 1, 1, 1])
# Train the perceptron
weights, bias = train_perceptron(X, y)
# Final predictions
print("Final predictions:")
for i in range(len(X)):
    print(f"Input: {X[i]}, Output: {predict([X[i]], weights, bias)[0]}")

```

14 A) ANN

```

Import numpy as np
#creating the input array
X=np.array([[1,0,1,0],[1,0,1,1],[0,1,0,1]])
Print('\n Input:')
Print(X)
Y=np.array([[1],[1],[0]])
Print('\n Actual Output:')
Print(y)

```

```

Print"\n shape of Input:", X.shape)

def sigmoid(x):
    return 1/(1 + np.exp(-x))

derivative of Sigmoid Function
def derivatives_sigmoid(x):
    return x * (1 - x)

lr=0.1 #learning rate

inputlayer_neurons = X.shape[1] # no. of features in dataset
hiddenlayer_neurons = 3 # no. of hidden layers neurons
output_neurons = 1 # no. of neurons at output layer

#Initializing weight and bias
wh=np.random.uniform(size=(inputlayer_neurons, hiddenlayer_neurons))
bh= np.random.uniform(size=(1, hiddenlayer_neurons))
wout= np.random.unifor(size=(hiddenlayer_neurons, output_neurons))
bout= np.random.unifor(size=(1, output_neurons))

#training the model
For l in range(epoch):
    #Forward propagation
    hidden_layer_input1=np.dot(X,wh)
    hidden_layer_input=hidden_layer_input 1 +bh
    hidden_layer_activations=sigmoiud(hidden_layer_input)
    output_layer_input1=np.dot(hiddenlayer_activations, wout))
    output_layer_input= output_layer_input1bout
    output = sigmoid(output_layer_input)

    #Backpropagation
    E = y-output
    Slope_output_layer = derivatives_sigmoid(output)
    Slope_hidden_layer = derivatives_sigmoid(hiddenlayer_activations)
    d_output = E * slope_Output_layer
    Error_at_hidden_layer = d_output.dot(wout.T)
    d_hiddenlayer= Error_at_hidden_layer * slope_Hidden_layer

```

```

wout+=hiddenlayer_activations.T.dot(d_output) * lr
bout+=np.sum(d_output,axis=0, keepdims=True) * lr
wh+=X.T.dot(d_hiddenlayer) * lr
bh+=np.sum(d_hiddenlayer,axis=0, keepdims=True) * lr

```

15 A) RFE

```

import pandas as pd

from sklearn.datasets import load_Iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

#Load the Iris dataset
Iris_data = load_Iris()
X = Iris_data.data
Y = Iris_data.target

#Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Intitalize the Random Forest Classifier
random_forest = RandomForestClassifier(n_estimators=100, random_state = 42)

#Train the Model
random_forest.fit(X_train, y_train)

#Make Predictions on the test set
y_pred = random_forest.predict(X_test)

#Evaluate the Model
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_rep)
print("\nConfusion Matrix:\n", conf_matrix)

```

15B) Decision Tree Classification

Read the CSV Files

```
import pandas as pd
```

```
import numpy as np
```

```
df=pd.read_csv('PlayTennis.csv')
```

```
from sklearn.preprocessing import LabelEncoder
```

```
enc = LabelEncoder()
```

```
df_num_cat=pd.DataFrame()
```

```
df_num_cat['Outlook']= enc.fit_transform(df['Outlook'])
```

```
df_num_cat['Temperature']= enc.fit_transform(df['Temperature'])
```

```
df_num_cat['Humidity']= enc.fit_transform(df['Humidity'])
```

```
df_num_cat['Wind']= enc.fit_transform(df['Wind'])
```

```
df_num_cat['Play Tennis']= enc.fit_transform(df['Play Tennis'])
```

```
df_num_cat
```

```
x=df_num_cat.drop(['Play Tennis'],axis=1)
```

```
y=df_num_cat['Play Tennis']
```

Split dataset into training set and test set

```
from sklearn.model_selection import train_test_split # Import train_test_split function
```

```
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=1) #  
70% training
```

and 30% test

```
print(X_train)
```

```
print(X_test)
```

```
print(y_train)
```

```
print(y_test)
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
```

```
df_clf=DecisionTreeClassifier(criterion='entropy')
```

```
#df_clf=DecisionTreeClassifier()
```

```
df_clf.fit(X_train,y_train)
```



```

y_pred=df_clf.predict(X_test)
print(y_train)
print(y_test)
print(y_pred)
print("Accuracy")
print(metrics.accuracy_score(y_test,y_pred))

# Plot the decision tree
plt.figure(figsize=(12, 8))

plot_tree(df_clf, filled=True, feature_names=X_train.columns, class_names=df['Play
Tennis'].unique())

plt.show(

```

16A) BOOSTING ENABLE METHOD

```

import pandas as pd

from sklearn.datasets import load_Iris

from sklearn.model_selection import train_test_split

from sklearn.ensemble import AdaBoostClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

#Load the Iris dataset

Iris_data -= load_Iris()

X = Iris_data.data

Y = Iris_data.target

#Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Intitalize the AdaBoost Classifier with Decision Tree as base estimator

adaboost = AdaBoostClassifier(n_estimators=50, random_state = 42)

#Train the Model

adaboost.fit(X_train, y_train)

#Make Predictions on the test set

y_pred = adaboost.predict(X_test)

#Evaluate the Model

accuracy = accuracy_score(y_test, y_pred)

```

```
classification_rep = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_rep)
print("\nConfusion Matrix:\n", conf_matrix)
```